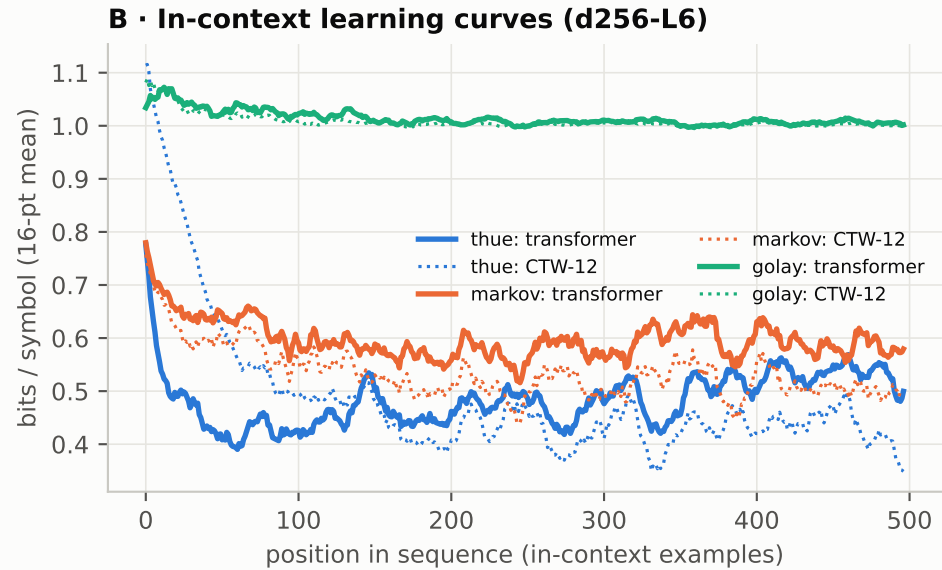
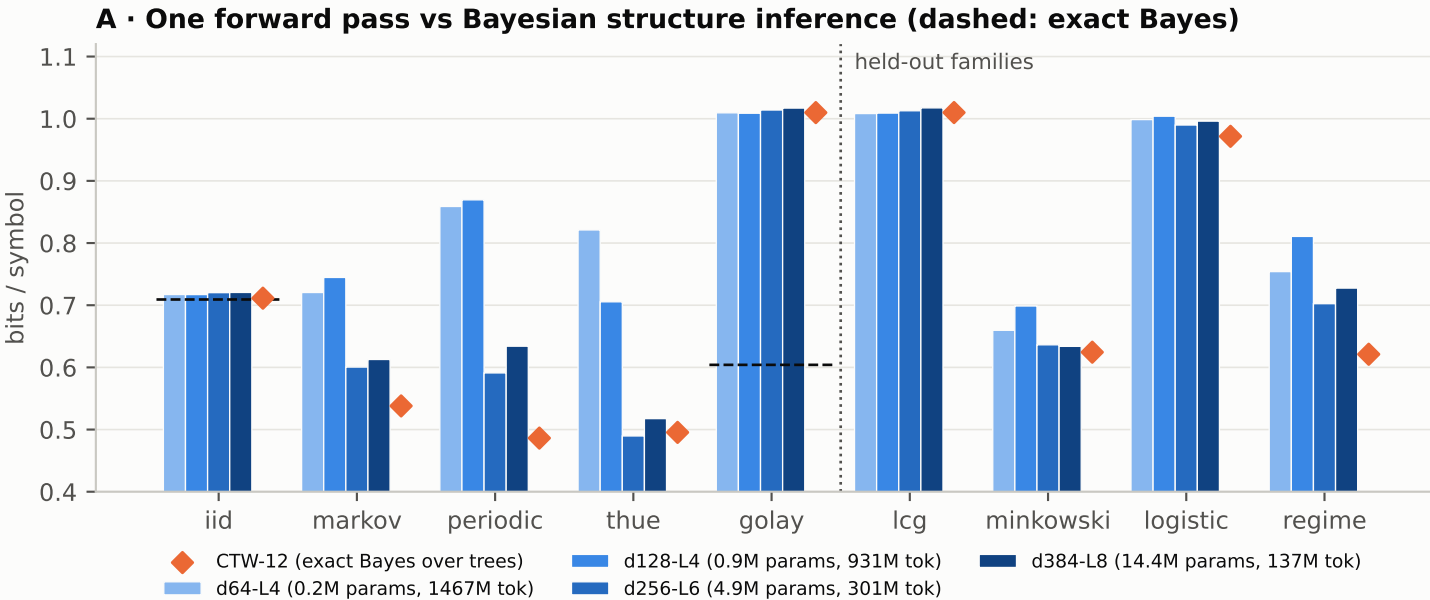


Fig. 9 · Amortised inference: a transformer meta-trained on a prior does structure inference in context



E007. Causal transformers meta-trained on a prior over generators (iid, Markov up to order 5, noisy periodic, noisy Thue-Morse, Golay), 220 s each on one H100 (compute-matched; about \$0.95 total). 64 eval sequences x 512 bits per family. MCB ≤ 0.02 bits on every family, held-out included. Golay stays near 1 bit (parity barrier) vs Bayes 0.60.